

Altitude-Matched Retrieval and Agentic Mixture-of-Experts for a Local-First Medical Reasoning Companion

Ayush Kumar
Healix Project
Independent Research
Email: preetndt@gmail.com

Abstract—Retrieval-augmented generation (RAG) grounds large language models (LLMs) in domain knowledge, but two problems remain when the users are patients rather than clinicians. Flat retrieval returns leaf passages that answer a specific question well, yet a broad question such as “how does chronic stress affect the body” has no single passage that answers it; the answer is spread across many. And a single dense embedding of a colloquial, under-specified query often fails to match a corpus written in clinical register. We describe Healix, an on-device medical companion, and a retrieval method we call *altitude-matched retrieval* that targets both problems together. A small language model drafts a hypothetical answer at two levels of abstraction in one call; the specific version joins the query to search the leaf passage index, and the broad version searches a RAPTOR-style summary tree. Coupling Hypothetical Document Embeddings (HyDE) with recursive summarization this way—matching each hypothetical to the index tier of the same granularity—does not appear in prior work. Healix also uses an agentic Mixture-of-Medical-Experts (MoME) router that gates queries to specialist retrieval lenses at the system level, and it runs entirely offline on one commodity GPU. We evaluate the stack with pre-registered hypotheses across four experiments: paraphrased and keyword query variation, subtopic coverage on broad questions, and end-to-end latency. The retrieval results are candidly negative—on this curated corpus a strong dense encoder is near ceiling, hybrid fusion and HyDE add no significant recall, and at a fixed evidence budget summary nodes cover fewer answer subtopics than leaf passages. The positive result is a systems one: moving the query-side HyDE call to a sub-billion-parameter model cuts end-to-end time-to-first-token from 14.7s to 6.4s (2.3 $\times$), far exceeding the isolated call’s speedup, because on a single shared GPU the cost of an auxiliary model is dominated by contention rather than its own size. We discuss the safety posture of a deliberately disclaimer-free assistant and release the system, benchmarks, and all artifacts as open source.

Index Terms—retrieval-augmented generation, dense retrieval, hierarchical retrieval, mixture of experts, medical question answering, on-device inference, large language models

I. INTRODUCTION

Large language models can produce fluent medical explanations but hallucinate facts they were not reliably trained on, and privacy concerns discourage sending personal health questions to hosted APIs. Retrieval-augmented generation (RAG) [1] mitigates the first problem by grounding generation in retrieved evidence; running the entire pipeline locally mitigates the second. Building such a system for lay medical

use, however, exposes two failure modes that standard RAG does not handle.

The altitude gap. Users ask questions at very different levels of abstraction. “What is the maximum daily dose of acetaminophen?” is answered by a single leaf passage. “How does chronic stress affect the body overall?” is not: no single passage contains the answer, which is an aggregate over cardiovascular, immune, digestive, and neurological sources. Flat dense retrieval, tuned to return the top passages, structurally cannot serve the second class. Recursive summarization indexes such as RAPTOR [3] add higher-level summary nodes to close this gap on the *index* side, but still search those nodes with the raw user query, whose embedding sits far from any summary.

The phrasing gap. Patients phrase questions colloquially and incompletely (“my chest feels tight when I’m stressed”), whereas the corpus is written in clinical register. Single-vector dense retrieval blurs exact terms and is sensitive to this mismatch. Hypothetical Document Embeddings (HyDE) [2] address this on the *query* side by embedding a generated hypothetical answer instead of the question, but classic HyDE produces one hypothetical against one flat index.

HyDE and RAPTOR act on the same axis—abstraction—from opposite ends, so they can be combined. We generate a hypothetical at two altitudes in one language-model call and send each to the tier whose granularity it matches: the specific hypothetical joins the query for leaf retrieval, and the broad one searches the summary tree. Because both come from a single call, the second altitude costs nothing extra.

We integrate altitude-matched retrieval into Healix, a local-first medical companion, alongside a system-level sparse Mixture-of-Medical-Experts router, a hybrid dense/sparse retriever with reranking, and a persona-controlled generator, all served offline from a single GPU. Our contributions are:

- 1) **Altitude-matched retrieval**, a novel coupling of HyDE and RAPTOR in which multi-granularity hypotheticals are routed to matching index tiers (Section IV).
- 2) **Agentic Mixture-of-Medical-Experts (MoME)**, an embedding-gated sparse router over specialist retrieval lenses that expands evidence coverage without a second generation pass (Section V).

- 3) **An honest, pre-registered evaluation** of the full stack on MedQuAD/PubMedQA—paraphrased and keyword query variation, broad-question subtopic coverage, and end-to-end latency—whose retrieval results are largely negative and reported as such, and whose positive finding is architectural: on a shared single GPU, auxiliary-model cost is governed by contention, and offloading the query-side model yields a $2.3\times$ end-to-end speedup (Sections VI–VII).

II. RELATED WORK

Dense and hybrid retrieval. Dense passage retrieval [4] maps queries and passages to a shared vector space; we use BGE embeddings [8] over a FAISS [9] index. Lexical BM25 [5] remains competitive for exact terms (drug names, rare conditions); fusing the two with Reciprocal Rank Fusion [6] and cross-encoder reranking [7] is now standard practice, and Anthropic’s Contextual Retrieval [10] shows that prepending situating context before indexing further improves recall. Healix uses all of these as its base retriever.

Query- and index-side augmentation. HyDE [2] embeds a generated hypothetical answer to bridge the query–corpus register gap. RAPTOR [3] builds a tree of recursive summaries so that broad queries can retrieve aggregated evidence. Both are single-altitude: HyDE produces one hypothetical; RAPTOR searches its tree with the raw query. We are not aware of prior work that generates hypotheticals at multiple abstraction levels and matches each to a corresponding index tier.

Mixture of experts and agents. Sparsely-gated MoE [11] and Switch Transformers [12] route tokens to expert subnetworks *inside* a model. More recently, mixture-of-agents approaches [13] route at the system level across LLM instances. Healix applies sparse gating one level up again—over medical *specialties*, each with its own retrieval lens—so that routing shapes evidence gathering rather than parameter selection.

Medical QA and safety. MedQuAD [14] and PubMedQA [15] are widely used medical QA corpora and form our knowledge base and evaluation source. Large medical LLMs such as Med-PaLM [16] target clinician-grade benchmarks; Healix instead targets lay users on-device, trading model scale for privacy, controllability, and grounded retrieval.

III. SYSTEM ARCHITECTURE

Healix processes each turn through the pipeline in Fig. 1 (described textually here). An input message first passes a *scope gate* that resolves greetings, identity, and clearly non-medical requests to canned responses with zero model cost, admitting only medical queries to the full pipeline. For an admitted query the system (i) runs the MoME router to select specialist lenses, (ii) generates two-altitude HyDE hypotheticals, (iii) retrieves leaf evidence via hybrid dense/sparse search under each expert lens and overview evidence from the RAPTOR tree, (iv) composes a single grounded prompt, and (v) streams a persona-controlled answer from a local LLM. Persistent conversation memory supplies continuity across turns. All components are served offline; the LLM runs on a

```

message → scope gate → [medical]
→ MoME router (specialist lenses)
→ two-altitude HyDE (one small-LM call)
→ leaf: hybrid dense+BM25+RRF+rerank
→ broad: RAPTOR summary-tree search
→ prompt compose → local LLM stream

```

Fig. 1. The Healix per-turn pipeline. Altitude-matched retrieval (Section IV) spans the two retrieval lines; the MoME router (Section V) conditions the leaf line.

local Ollama server and every index is memory-mapped from disk.

IV. ALTITUDE-MATCHED RETRIEVAL

A. Hierarchical index

The leaf tier L_0 is the corpus of $N=239,644$ passage chunks embedded with BGE and indexed in FAISS using 8-bit scalar quantization, chosen because it preserves absolute cosine scores (mean score error 1.7×10^{-4} versus a flat index) so that score thresholds remain valid. We construct two summary tiers offline. Leaf embeddings are reconstructed directly from the quantized index (no re-embedding) and clustered with k -means into $K_1=1500$ topic clusters; the local LLM summarizes each cluster’s representative members into a titled overview paragraph, forming tier L_1 . The L_1 summaries are themselves embedded and clustered into $K_2=120$ domain nodes (L_2). The build is resumable and processes clusters largest-first so partial builds cover the highest-mass topics. The resulting tree contains $1500+114=1614$ overview nodes.

B. Two-altitude hypotheticals

Given query q , a single small-LM call emits two hypotheticals: a *specific* sentence h_s carrying concrete clinical detail, and a *broad* sentence h_b naming the overarching topic. We deliberately use a sub-billion-parameter model here because the hypothetical is consumed only as an embedding target, not shown to the user; a lenient parser falls back to the raw generation when a small model ignores the output format, so the mechanism never silently no-ops.

C. Matched routing

Leaf retrieval embeds the concatenation $[q; h_s]$ and searches L_0 through the full hybrid pipeline. Overview retrieval embeds h_b (falling back to q) and searches $L_1 \cup L_2$; the top overview nodes are *prepended* to the evidence block so the generator reads the frame before the details. Formally, for tier L_t with granularity g_t , we search with the hypothetical h whose target granularity is closest to g_t :

$$\text{evidence} = \bigcup_t \text{TopK}(L_t, \phi(h_{\arg \min_a |g_a - g_t|})),$$

where ϕ is the shared embedding function. Classic HyDE is the special case of a single tier and single hypothetical; classic RAPTOR is the special case of searching all tiers with $\phi(q)$. Matching the hypothetical’s altitude to the tier aligns query and corpus geometry on both axes at once, and because both hypotheticals come from one call, the second altitude is

free. The mechanism is *fail-open*: any error, a missing tree, or an unavailable LM degrades gracefully to standard hybrid retrieval on q .

V. AGENTIC MIXTURE-OF-MEDICAL-EXPERTS

Rather than route tokens (as in intra-model MoE) or whole LLM instances (as in mixture-of-agents), Healix routes over medical *specialties*. Each expert e carries (i) a semantic prototype p_e encoded once, (ii) a keyword prior, and (iii) a retrieval *lens*—a query reformulation template. The gate computes cosine similarity between the query embedding and each prototype, adds keyword and imaging-modality boosts, and applies a sharpened softmax; the top- k experts above a floor are activated:

$$w_e = \text{softmax}_\tau(\langle \phi(q), p_e \rangle + \beta \mathcal{K}[\text{kw}] + \gamma \mathcal{K}[\text{modality}]).$$

Each activated expert reformulates q through its lens and gathers specialty-specific evidence; the union grounds a *single* synthesis pass framed as an integrated panel, so coverage improves without the cost of multiple generations. The current roster covers Cardiology, Pulmonology, Neurology, Dermatology, Gastroenterology, Musculoskeletal, Pharmacology, Psychophysiology, and a default General expert. Activated experts and gate weights are surfaced to the user as an explainability signal.

VI. IMPLEMENTATION AND LATENCY ENGINEERING

Healix is served by a FastAPI backend streaming tokens over Server-Sent Events to a browser UI; generation uses a local Ollama server (default `qwen2.5:7b-instruct`). The knowledge index (175 MB quantized) and chunk store are memory-mapped. Because the multi-stage pipeline runs several operations *before* the first generated token, query-side latency governs perceived responsiveness, and HyDE is the largest such operation—it invokes a language model on the critical path. Since the hypothetical is consumed only as an embedding target, we run HyDE on a sub-billion-parameter model rather than the generation model.

Two measurements quantify the effect (Table I), and their difference is instructive. In isolation, the HyDE call itself shrinks only modestly, from 3.76 s to 3.01 s (mean over 25 queries; paired bootstrap 95% CI on the gap [0.67, 0.85] s). End to end on the warm server, however, time-to-first-token falls from 14.7 s to 6.4 s and total latency from 16.6 s to 8.7 s (mean over 10 queries each)—a $2.3\times$ reduction far larger than the isolated saving. The gap is GPU contention: on a single 8 GB device the 7B HyDE call competes with the embedding model and cross-encoder reranker for compute, so it runs much slower inside the live pipeline than alone. Offloading it to a 0.5B model removes that contention. The lesson generalizes to any multi-stage local pipeline that time-shares one accelerator: the cost of an auxiliary model is dominated by contention, not by its own size. All stages remain fail-open, so any component can be disabled or fail without breaking a response.

TABLE I
HYDE LATENCY, WARM, SINGLE 8 GB (RTX-4060-CLASS) GPU.
ISOLATED = THE HYDE CALL ALONE; END-TO-END = FULL PIPELINE
TIME-TO-FIRST-TOKEN AND TOTAL.

Measurement	HyDE on 7B	HyDE on 0.5B
Isolated HyDE call (s)	3.76	3.01
End-to-end TTFT (s)	14.7	6.4
End-to-end total (s)	16.6	8.7

VII. EVALUATION

We ran four experiments, each with its hypothesis stated before execution, and we report every outcome including the failures. All harnesses, seeds, and result files are released.

A. Protocol

Because the corpus questions appear verbatim in the index, evaluating retrieval with the raw question is circular (dense retrieval trivially self-matches). We therefore use a *query-variation* protocol: we sample n question–chunk pairs, paraphrase each question into colloquial patient phrasing with an LLM, keep the *original chunk* as the gold target, and measure whether each retrieval configuration recovers it in the top k . This isolates robustness to phrasing, exactly the property hybrid and HyDE claim to improve. We report Recall@1/5/10 and MRR@10 for three configurations: **dense** (FAISS + cross-encoder rerank), **hybrid** (dense + contextual BM25 + RRF + rerank), and **hybrid+HyDE** (hybrid search over $[q; h_s]$). Sample size n , seed, and paraphrase model are fixed and released with the code.

B. Retrieval results

Table II gives the benchmark, and the outcome is negative: the three configurations differ by at most two points on any metric, share an identical Recall@10 of 0.880, and HyDE spends about three seconds per query without improving recall. We report the full result rather than a favorable slice. The parity has three likely causes. The cross-encoder reranker is strong enough to equalize the candidate sets that dense and hybrid produce, which is the expected behavior once reranking, not first-stage recall, decides the ordering. An LLM paraphrase of an in-domain question also keeps most of its medical content, so the dense encoder rarely misses by enough for lexical or hypothetical augmentation to help. And the metric only credits recovery of one gold chunk, so it cannot reward the aggregate-evidence case the hierarchical arm is built for. The engineering consequence is clear: for in-domain single-target retrieval, the cost of HyDE is not repaid in recall, which is why we keep it cheap and optional rather than mandatory. Showing a win for hybrid or HyDE would need harder queries than this benchmark contains—out-of-distribution phrasing or rare drug and gene names—and building that set is the main evaluation gap we leave open.

C. Keyword queries

Dense encoders are trained on natural language and are expected to degrade on terse keyword input, where BM25

TABLE II
QUERY-VARIATION RETRIEVAL ON PARAPHRASED
MEDQUAD/PUBMEDQA ($n=150$, $k=10$, SEED 13, 100%
PARAPHRASED). HIGHER IS BETTER; LAST COLUMN IS MEAN QUERY-SIDE
LATENCY.

Configuration	R@1	R@5	R@10	MRR@10	Lat. (s)
Dense	0.700	0.853	0.880	0.760	0.14
Hybrid	0.700	0.820	0.880	0.755	0.19
Hybrid + HyDE	0.667	0.847	0.880	0.739	3.24

should help. We repeated the protocol with deterministic keyword queries (content words only, e.g. “glaucoma treatment”) over $n=500$ samples. The augmentations again did not help materially: hybrid first-stage Recall@10 was 0.964 versus 0.958 for dense, a gap of 0.006 with a bootstrap 95% CI of $[0.000, 0.014]$ that includes zero. Hybrid rescued three queries that dense missed and lost none, so it never hurt, but the effect is not significant. Reranking added a small, consistent precision gain (dense Recall@1 $0.916 \rightarrow 0.930$; hybrid $0.914 \rightarrow 0.936$). The BGE encoder handles keyword queries well enough on this corpus that lexical fusion has little to recover.

D. Broad questions and the hierarchical arm

We tested altitude matching directly on the regime it targets. We authored 24 broad questions, each with an objective set of textbook subtopics (e.g., complications of diabetes: retinopathy, nephropathy, neuropathy, cardiovascular, foot, wound healing), and measured *subtopic coverage* at a fixed 2000-character evidence budget—counting how many subtopics appear in the retrieved text, which controls for summary nodes being longer than leaf passages. The result is negative and significant: flat leaf retrieval covered 0.67 of subtopics on average, the RAPTOR overview tier only 0.34, and their union 0.33 (overview minus leaf gap -0.33 , 95% CI $[-0.46, -0.20]$). Inspection shows why: cluster summaries generalize—“a range of complications”—and drop the specific terms that leaf passages state outright, so on a corpus that already contains dedicated passages per subtopic, summarization removes rather than adds coverage. Altitude matching does not help here, and prepending overview nodes can displace more informative leaf evidence. Where a summary tier would help is a corpus without per-subtopic leaf coverage; establishing that is future work.

E. System behavior

Under a 5-way concurrent load the server completed all requests with isolated conversation state and no errors, serializing generation behind a single lock. Adversarial and out-of-domain inputs (prompt-injection attempts, non-medical requests, empty and malformed messages) were refused or handled without leaking system instructions.

VIII. DISCUSSION AND LIMITATIONS

Safety posture. Healix is deliberately disclaimer-free and avoids ritual “consult a professional” closings, favoring spe-

cific, calm guidance. This design choice trades a conventional safety net for readability and must be revisited for any real deployment, particularly the handling of red-flag emergency presentations, where suppressing urgent guidance would be harmful. **Evaluation scope.** Our benchmark measures retrieval, not answer quality; a human or LLM-judge study of grounded answer accuracy is the most important next step. **Model ceiling.** The 7B generator bounds clinical nuance; the architecture is model-agnostic and benefits directly from stronger local models. **Paraphrase confound.** The query-variation protocol paraphrases questions with an LLM. The retriever’s encoder is a separate embedding model, so this is less circular than an LLM-as-judge study, but paraphrase artifacts remain; a human-paraphrased query set would be stronger. We fix the seed and release all artifacts for reproduction.

IX. CONCLUSION

We described Healix and altitude-matched retrieval, which couples HyDE and RAPTOR by drafting hypotheticals at two granularities and routing each to the matching tier of a hierarchical index, alongside a system-level Mixture-of-Medical-Experts router. We then tested the stack the hard way, with pre-registered hypotheses, and most of them failed: on a curated medical QA corpus a strong dense encoder leaves essentially no recall headroom for hybrid fusion or HyDE, and summary nodes cover fewer answer subtopics than the leaf passages they abstract. We take these results seriously rather than explaining them away. They suggest a practical rule for builders: measure whether the corpus already gives dense retrieval per-subtopic leaf coverage before paying for augmentation machinery. The durable positive finding is architectural—in a multi-stage pipeline time-sharing one accelerator, auxiliary model cost is set by contention, not parameter count, and offloading the query-side model yielded a $2.3\times$ end-to-end speedup. Whether altitude matching pays off on corpora *without* dense leaf coverage remains open, and our released benchmarks make that test straightforward to run. Code and artifacts are public for reproduction.

REFERENCES

- [1] P. Lewis *et al.*, “Retrieval-augmented generation for knowledge-intensive NLP tasks,” in *Proc. NeurIPS*, 2020.
- [2] L. Gao, X. Ma, J. Lin, and J. Callan, “Precise zero-shot dense retrieval without relevance labels,” in *Proc. ACL*, 2023.
- [3] P. Sarthi *et al.*, “RAPTOR: Recursive abstractive processing for tree-organized retrieval,” in *Proc. ICLR*, 2024.
- [4] V. Karpukhin *et al.*, “Dense passage retrieval for open-domain question answering,” in *Proc. EMNLP*, 2020.
- [5] S. Robertson and H. Zaragoza, “The probabilistic relevance framework: BM25 and beyond,” *Found. Trends Inf. Retr.*, 2009.
- [6] G. V. Cormack, C. L. A. Clarke, and S. Buettcher, “Reciprocal rank fusion outperforms Condorcet and individual rank learning methods,” in *Proc. SIGIR*, 2009.
- [7] N. Reimers and I. Gurevych, “Sentence-BERT: Sentence embeddings using Siamese BERT-networks,” in *Proc. EMNLP*, 2019.
- [8] S. Xiao *et al.*, “C-Pack: Packed resources for general Chinese embeddings,” in *Proc. SIGIR*, 2024.
- [9] J. Johnson, M. Douze, and H. Jégou, “Billion-scale similarity search with GPUs,” *IEEE Trans. Big Data*, 2019.
- [10] Anthropic, “Introducing contextual retrieval,” Tech. Report, 2024.

- [11] N. Shazeer *et al.*, “Outrageously large neural networks: The sparsely-gated mixture-of-experts layer,” in *Proc. ICLR*, 2017.
- [12] W. Fedus, B. Zoph, and N. Shazeer, “Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity,” *J. Mach. Learn. Res.*, 2022.
- [13] J. Wang *et al.*, “Mixture-of-agents enhances large language model capabilities,” *arXiv:2406.04692*, 2024.
- [14] A. Ben Abacha and D. Demner-Fushman, “A question-entailment approach to question answering,” *BMC Bioinformatics*, 2019.
- [15] Q. Jin *et al.*, “PubMedQA: A dataset for biomedical research question answering,” in *Proc. EMNLP*, 2019.
- [16] K. Singhal *et al.*, “Large language models encode clinical knowledge,” *Nature*, 2023.